

Generalized Linear Models

Applications

Generalized linear models (GLMs)

Generalized Linear Models provide a unified framework for modeling outcomes from the various family of distributions, including:

- Continuous outcomes (Gaussian → linear regression)
- Binary outcomes (Binomial → logistic regression)
- Count outcomes (Poisson → rate models)

Core idea

All GLMs share the same structure:

$$g(\mu_i) = X_i^\top \beta$$

Where:

- Y_i follows a distribution from the **exponential family**
- $\mu_i = E(Y_i | X_i)$
- $g(\cdot)$ is a **link function**
- $X_i^\top \beta$ is the linear predictor

Three components of a GLM

1. Random component

Distribution of Y (Normal, Binomial, Poisson)

2. Systematic component

Linear predictor $X^\top \beta$

3. Link function

Connects mean outcome to predictors

- Identity \rightarrow Linear regression
- Logit \rightarrow Logistic regression
- Log \rightarrow Poisson regression

Why GLMs matter in health research

They allow us to model:

- Mean differences (QALY)
- Odds ratios (disease risk)
- Incidence rate ratios (person-time data)

Within one unified framework.

Outcome Type	Distribution	Link	Interpretation
Continuous	Gaussian	Identity	Mean difference
Binary	Binomial	Logit	Odds ratio
Count	Poisson	Log	Rate ratio

Linear regression: Ketamine for Chronic Pain (QALY)

Clinical Question

Does ketamine dosage and patient characteristics influence **Quality-Adjusted Life Years (QALY)?**

We model the expected outcome:

$$E(Y | X)$$

Where:

- ($Y = QALY$)
- ($X =$) treatment characteristics + patient covariates

Model specification

Linear regression assumes:

$$Y_i = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_p X_{ip} + \varepsilon_i$$

with:

$$\varepsilon_i \sim \mathcal{N}(0, \sigma^2)$$

Thus:

$$E(Y_i | X_i) = X_i^\top \beta$$

$$Var(Y_i | X_i) = \sigma^2$$

Interpretation of coefficients

For a continuous predictor (X_j):

$$\beta_j = \frac{\partial E(Y)}{\partial X_j}$$

- Expected change in QALY for a 1-unit increase in X_j , holding other variables constant.

For categorical predictors:

β_j = Difference in mean QALY vs reference group

Estimation principle

Parameters are estimated via **Ordinary Least Squares (OLS)**:

$$\hat{\beta} = (X^\top X)^{-1} X^\top Y$$

OLS minimizes:

$$\sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

Assumptions

- Linearity
- Independence
- Homoscedasticity
- Normal residuals
- No strong multicollinearity

Questions

1. Load and inspect the `ketapain.csv` Dataset.

```
keta <- read.csv("datasets/ketapain.csv")
```

```
head(keta)
```

```
  patientID sexe age av_dose level_dose cum_dose cum_days perfusion cost
1       6317 female  35   0.700  low dose    269.5      5     24 4689.95
2       2517 female  76   0.700  low dose    269.5      5     24 4689.95
3      1023 female  18   0.075  low dose     63.0     12     4 11165.30
4      3002 female  50  74.000  low dose    222.0      3     24 2829.20
5       7702 male   46  74.000  low dose    222.0      3     24 2829.20
6       439 female  54   0.500  low dose   105.0      3     6 2817.50
  qaly mode
1 0.32938894 continu
2 0.69094614 continu
3 0.31734359 discontinu
4 0.04861697 continu
5 0.63407075 continu
6 0.43630836 discontinu
```

```
str(keta)
```

```
'data.frame': 184 obs. of 11 variables:
 $ patientID : int  6317 2517 1023 3002 7702 439 1224 1616 2924 1317 ...
 $ sexe       : chr "female" "female" "female" "female" ...
 $ age        : int  35 76 18 50 46 54 51 50 43 30 ...
 $ av_dose    : num  0.7 0.7 0.075 74 74 0.5 100 1.6 100 0.7 ...
 $ level_dose: chr "low dose" "low dose" "low dose" "low dose" ...
 $ cum_dose   : num  270 270 63 222 222 ...
 $ cum_days   : int  5 5 12 3 3 3 3 5 3 5 ...
 $ perfusion  : num  24 24 4 24 24 6 3 24 3 24 ...
 $ cost       : num  4690 4690 11165 2829 2829 ...
 $ qaly       : num  0.3294 0.6909 0.3173 0.0486 0.6341 ...
 $ mode       : chr "continu" "continu" "discontinu" "continu" ...
```

Before modeling, check:

- Is `qaly` continuous?
- Any missing values?
- Correct variable types?

2. Perform some basic cleaning.

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

- Why convert to factors?

Because categorical variables must be treated as group comparisons, not numeric scales.

3. Fit the linear regression.

We estimate:

$$E(\text{QALY} | X) = \beta_0 + \beta_1 \text{Age} + \beta_2 \text{Sex} + \beta_3 \text{Average Dose} + \dots$$

Call:

```
lm(formula = qaly ~ age + sexe + level_dose + cum_dose + cum_days +  
    perfusion + mode, data = keta)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.4744	-0.1313	0.0092	0.1148	0.4621

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.3580489	0.1467429	2.440	0.0157 *
age	0.0013680	0.0012007	1.139	0.2561
sexemale	0.0306974	0.0350236	0.876	0.3820
level_dose	-0.0529579	0.1371372	-0.386	0.6998
level_dose	0.0373005	0.1404256	0.266	0.7908
cum_dose	-0.0003105	0.0001972	-1.575	0.1171
cum_days	-0.0030772	0.0074231	-0.415	0.6790
perfusion	0.0045027	0.0027079	1.663	0.0981 .

```

modediscontinu      0.0627110  0.0582525   1.077   0.2832
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1922 on 175 degrees of freedom
Multiple R-squared:  0.04889,  Adjusted R-squared:  0.005413
F-statistic: 1.124 on 8 and 175 DF,  p-value: 0.349

```

4. Interpret key coefficients

	(Intercept)	age	sexemale
	0.3580489014	0.0013679681	0.0306973972
level_doselow	dose	level_dose	midium
-0.0529579342		0.0373004898	level
cum_days		perfusion	
-0.0030772233		0.0045027387	modediscontinu
			0.0627109957

5. Check model assumptions

Residuals:

$$\hat{\varepsilon}_i = Y_i - \hat{Y}_i$$

We check:

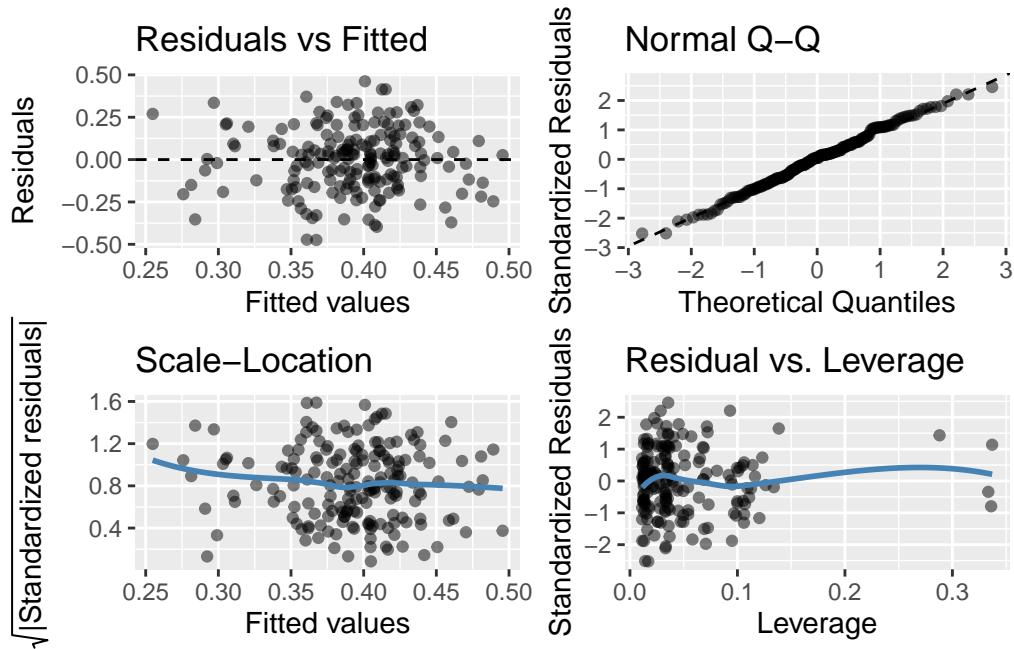
- Residual vs fitted → Linearity and homoscedasticity
- QQ-plot → Normality
- Influence diagnostics

Loading required package: ggplot2

```

Warning: `fortify(<lm>)` was deprecated in ggplot2 4.0.0.
i Please use `broom::augment(<lm>)` instead.
i The deprecated feature was likely used in the ggplot2 package.
Please report the issue at <https://github.com/tidyverse/ggplot2/issues>.

```



Interpretation of diagnostics

- Funnel shape → heteroscedasticity
 - Systematic curve → non-linearity
 - Heavy tails in QQ → non-normal residuals
 - Extreme leverage → influential observations
6. Predict QALY for a new patient

Theoretical prediction:

$$\hat{Y}_{new} = x_{new}^\top \hat{\beta}$$

Prediction interval:

$$\hat{Y}_{new} \pm t * n - p \sqrt{\hat{\sigma}^2 (1 + x_{new}^\top (X^\top X)^{-1} x_{new})}$$

Difference:

- Confidence interval → mean response
- Prediction interval → individual patient
- Helper function for predicting qaly:

```

predict_qaly <- function(age, sexe, av_dose, level_dose,
                         cum_dose, cum_days, perfusion, mode) {

  newdata <- data.frame(
    age = age,
    sexe = factor(sexe, levels = levels(keta$sexe)),
    av_dose = av_dose,
    level_dose = factor(level_dose, levels = levels(keta$level_dose)),
    cum_dose = cum_dose,
    cum_days = cum_days,
    perfusion = perfusion,
    mode = factor(mode, levels = levels(keta$mode))
  )

  predict(fit_lm, newdata = newdata, interval = "prediction")
}

```

{r} #| echo: false predict_qaly(age = 50, sexe = “female”, av_dose = 0.7, level_dose = “low dose”, cum_dose = 269.5, cum_days = 5, perfusion = 24, mode = “continu”)

Logistic regression: Arthritis risk factors (arthritis)

Clinical Question

Which patient factors are associated with **arthritis** in a cohort dataset?

Outcome (binary):

$$Y = \begin{cases} 1 & \text{arthritis present} \\ 0 & \text{arthritis absent} \end{cases}$$

We want to model the **risk**:

$$P(Y = 1 | X)$$

with covariates age, gender, BMI, diabetes, smoking.

Why Logistic Regression?

A linear model can predict values outside ([0,1]). Logistic regression ensures predicted risks are valid probabilities by using a **link function**:

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

with:

$$p_i = P(Y_i = 1 | X_i)$$

Model Specification

$$\log\left(\frac{P(Y_i = 1 | X_i)}{1 - P(Y_i = 1 | X_i)}\right) = \beta_0 + \beta_1 X_{i1} + \cdots + \beta_p X_{ip}$$

Equivalently, the predicted probability is:

$$p_i = \frac{1}{1 + \exp(-X_i^\top \beta)}$$

Interpretation of coefficients

Coefficients are in **log-odds**:

β_j = change in log-odds per 1-unit increase in X_j

Exponentiating gives the **Odds Ratio (OR)**:

$$OR_j = e^{\beta_j}$$

Interpretation:

- (OR>1): increases odds (higher risk)
- (OR<1): decreases odds (protective association)
- (OR=1): no association

For categorical predictors, ORs are relative to the **reference category**.

Estimation principle

Parameters are estimated by **Maximum Likelihood Estimation (MLE)**.

Likelihood:

$$L(\beta) = \prod_{i=1}^n p_i^{y_i} (1 - p_i)^{1-y_i}$$

We typically use the log-likelihood and optimize numerically.

Inference

We test:

$$H_0 : \beta_j = 0$$

using Wald tests (z-statistics) or likelihood-based tests.

Confidence intervals for ORs are:

$$\exp(\hat{\beta}_j \pm 1.96 \cdot SE(\hat{\beta}_j))$$

Marginal Effects (Clinical interpretation)

Because ORs can be hard to interpret clinically, marginal effects translate to **probability changes**:

$$\frac{\partial p}{\partial x_j} = \beta_j, p(1-p)$$

So the probability impact depends on baseline risk p . That's why we often compute **average marginal effects**.

1. Load and inspect the dataset.

```
arthritis <- read.csv("datasets/arthritis.csv")
head(arthritis)
```

```

      id status heart.attack.relative gender age   bmi diabetes alcohol smoke
1 41475   Yes                      No Female 62 58.04      No      No  No
2 41477   Yes                      No Male 71 30.05     Yes      No Yes
3 41479   No                       No Male 52 27.56      No     Yes  No
4 41481   No                       No Male 21 23.34      No     Yes  No
5 41482   No                       No Male 64 33.64      No     Yes Yes
6 41483   Yes                      No Male 66 44.06     Yes      No  No

  prehypertension vegetarian covered.health
1                 No        No        No
2                 No        No       Yes
3                 No        No        No
4                 No        No        No
5                 No        No       Yes
6                 No        No       Yes

```

```
str(arthritis)
```

```

'data.frame': 4856 obs. of 12 variables:
 $ id           : int 41475 41477 41479 41481 41482 41483 41485 41486 41487 41489 ...
 $ status        : chr "Yes" "Yes" "No" "No" ...
 $ heart.attack.relative: chr "No" "No" "No" "No" ...
 $ gender        : chr "Female" "Male" "Male" "Male" ...
 $ age           : int 62 71 52 21 64 66 30 61 27 40 ...
 $ bmi           : num 58 30.1 27.6 23.3 33.6 ...
 $ diabetes       : chr "No" "Yes" "No" "No" ...
 $ alcohol        : chr "No" "No" "Yes" "Yes" ...
 $ smoke          : chr "No" "Yes" "No" "No" ...
 $ prehypertension: chr "No" "No" "No" "No" ...
 $ vegetarian     : chr "No" "No" "No" "No" ...
 $ covered.health : chr "No" "Yes" "No" "No" ...

```

2. Prepare the data.

Convert categorical variables to factors, and set the reference outcome category.

Why reference matters?

- It determines which group is the baseline for interpretation.
- Here we model the odds of **status = “Yes”** relative to **“No”** (depending on coding; see note below).

3. Fit a logistic regression model.

Predictors:

- age
- gender
- bmi
- diabetes
- smoke

Call:

```
glm(formula = status ~ age + gender + bmi + diabetes + smoke,  
     family = binomial, data = arthritis)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-5.702551	0.235275	-24.238	< 2e-16 ***
age	0.060980	0.002369	25.746	< 2e-16 ***
genderMale	-0.602215	0.074183	-8.118	4.74e-16 ***
bmi	0.050995	0.005474	9.315	< 2e-16 ***
diabetesYes	0.221479	0.100630	2.201	0.0277 *
smokeYes	0.551452	0.073706	7.482	7.33e-14 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for binomial family taken to be 1)

Null deviance: 5861.9 on 4855 degrees of freedom
Residual deviance: 4754.4 on 4850 degrees of freedom
AIC: 4766.4

Number of Fisher Scoring iterations: 5

4. Compute Odds Ratios (OR)

$$OR = e^\beta$$

(Intercept)	age	genderMale	bmi	diabetesYes	smokeYes
0.003337442	1.062877476	0.547597410	1.052317555	1.247921256	1.735772165

5. Which variables increase the odds of arthritis?

Rule:

- OR > 1 increases odds
- OR < 1 decreases odds

```
(Intercept)      age   genderMale      bmi  diabetesYes   smokeYes
0.003337442 1.062877476 0.547597410 1.052317555 1.247921256 1.735772165
```

6. Compute 95% confidence intervals (OR scale).

Waiting for profiling to be done...

	2.5 %	97.5 %
(Intercept)	0.002093571	0.005266448
age	1.058002639	1.067873530
genderMale	0.473264698	0.633017443
bmi	1.041106201	1.063698236
diabetesYes	1.024189001	1.519654333
smokeYes	1.502786941	2.006308314

- If the CI includes 1 → evidence of association is weak (at 5% level)
- If CI entirely > 1 → increased odds
- If CI entirely < 1 → decreased odds

7. Compute predicted probability of arthritis.

Predicted risk:

$$\hat{p}_i = P(Y_i = 1 | X_i)$$

```
[1] 0.73845319 0.58179137 0.15079310 0.02116597 0.46626765 0.54691630
```

8. Visualize predicted probability vs age.

```
library(ggplot2)

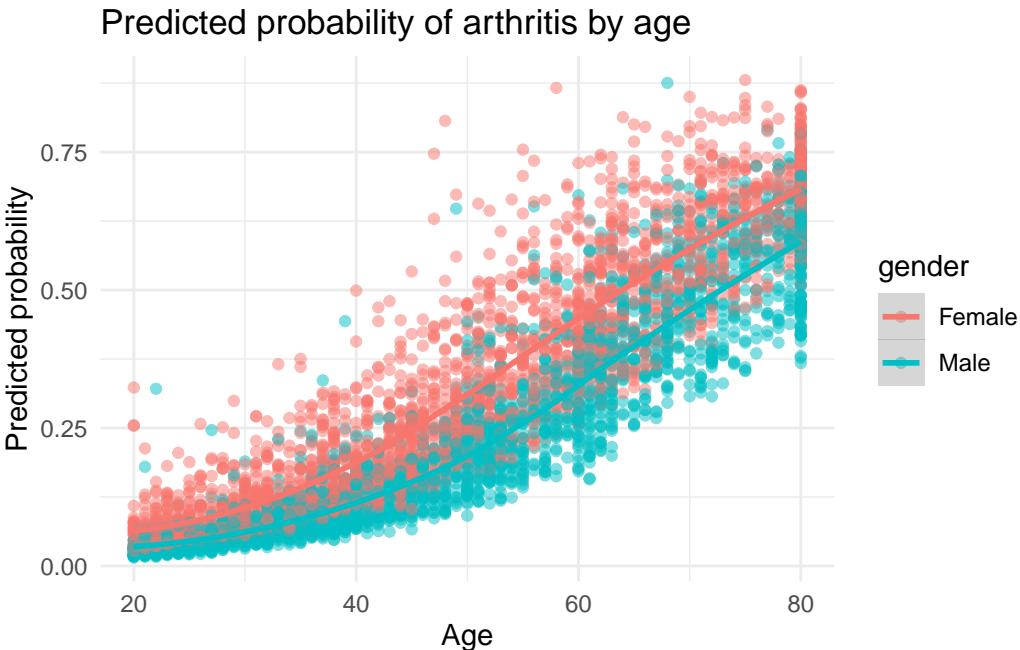
ggplot(arthritis, aes(age, pred_prob, color = gender)) +
  geom_point(alpha = 0.5) +
  geom_smooth(method = "loess") +
  labs(
    title = "Predicted probability of arthritis by age",
    x = "Age",
```

```

y = "Predicted probability"
) +
theme_minimal()

```

```
`geom_smooth()` using formula = 'y ~ x'
```



Teaching note:

- This is not a causal curve; it's model-based prediction given observed covariates.
- LOESS here is just for visualization, not the fitted logistic mean function.

9. Interpretation Questions (From ORs)

- Does age increase risk? ($OR_{age} > 1$?)
- Does BMI increase risk?
- Does smoking increase risk?
- Does gender affect risk?

10. Marginal effects of BMI and Age.

Coefficients are in log-odds so marginal effects express average probability change.

```

library(marginaleffects)

mfx <- avg_slopes(fit)
mfx

```

Term	Contrast	Estimate	Std. Error	z	Pr(> z)	S	2.5 %
age	dY/dX	0.00986	0.000282	34.95	<0.001	886.6	0.00931
bmi	dY/dX	0.00824	0.000859	9.60	<0.001	70.0	0.00656
diabetes	Yes - No	0.03677	0.017112	2.15	0.0317	5.0	0.00323
gender	Male - Female	-0.09738	0.011770	-8.27	<0.001	52.8	-0.12045
smoke	Yes - No	0.08958	0.011846	7.56	<0.001	44.5	0.06636
	97.5 %						
	0.01041						
	0.00993						
	0.07030						
	-0.07431						
	0.11280						

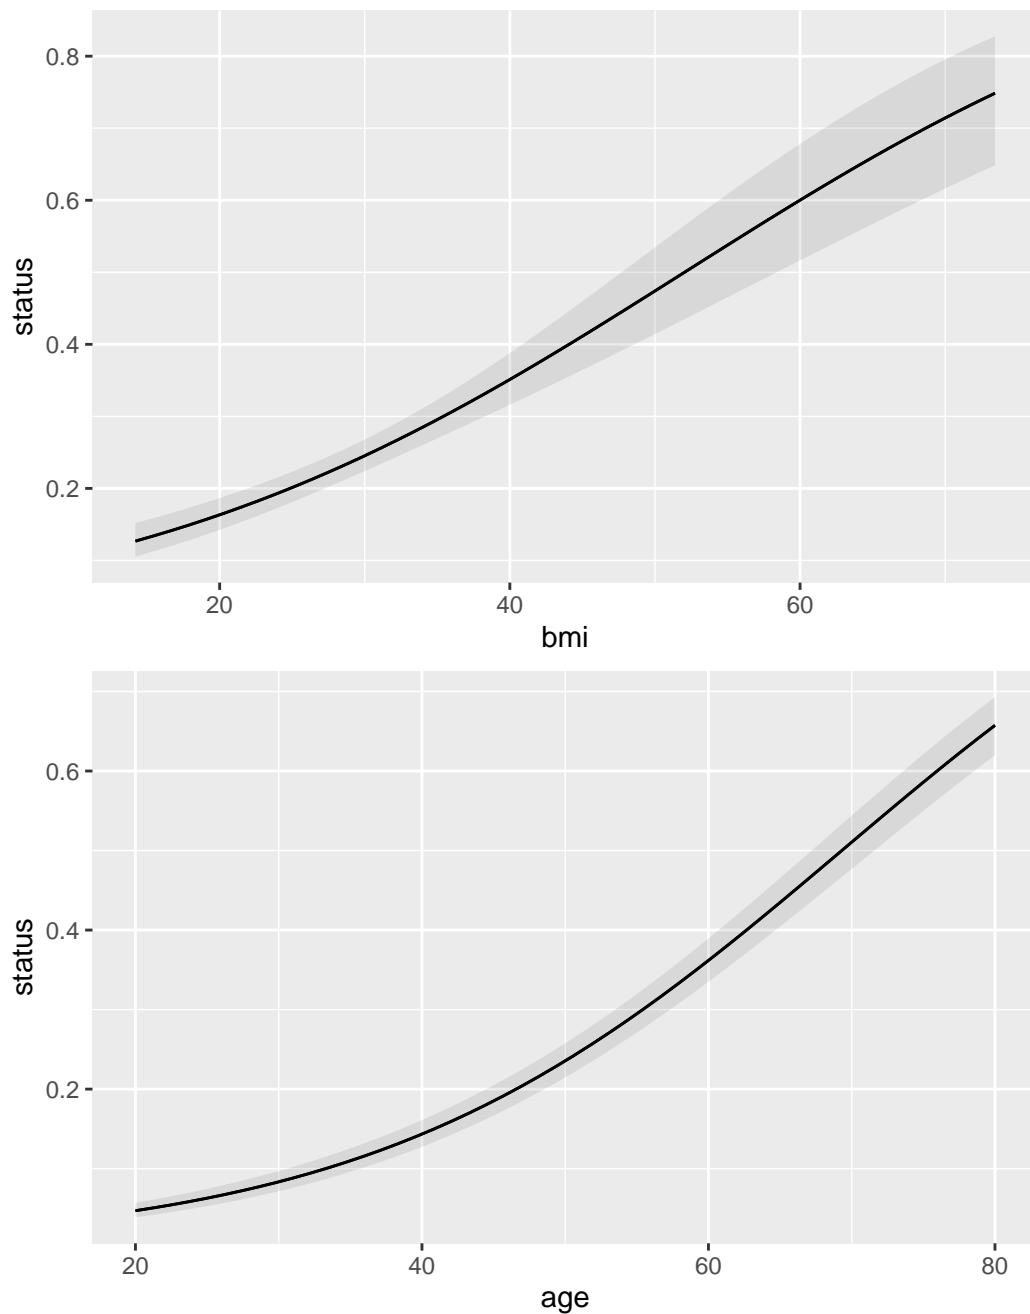
Type: response

- Marginal effects for BMI and age only

Term	Estimate	Std. Error	z	Pr(> z)	S	2.5 %	97.5 %
age	0.00986	0.000282	34.9	<0.001	886.6	0.00931	0.01041
bmi	0.00824	0.000859	9.6	<0.001	70.0	0.00656	0.00993

Type: response

Comparison: dY/dX



Interpretation:

- positive effect → increases predicted risk (on average)
- negative effect → decreases predicted risk

11. Predicting risk for new patients.

We now predict arthritis risk for specific individuals.

- Helper function

```
predict_risk <- function(age,
                           bmi,
                           gender = "Female",
                           diabetes = "No",
                           smoke = "No") {

  newdata <- data.frame(
    age = age,
    bmi = bmi,
    gender = factor(gender, levels = levels(arthritis$gender)),
    diabetes = factor(diabetes, levels = levels(arthritis$diabetes)),
    smoke = factor(smoke, levels = levels(arthritis$smoke)))
}

predict(
  fit,
  newdata = newdata,
  type = "response"
)
}
```

- Example: low-risk vs high-risk profiles

Low-risk:

- Female, Age 30, BMI 22, non-smoker, no diabetes

High-risk:

- Male, Age 70, BMI 35, smoker, diabetic

```
new_patients <- data.frame(
  age = c(30, 70),
  bmi = c(22, 35),
  gender = factor(c("Female", "Male"),
                  levels = levels(arthritis$gender)),
  diabetes = factor(c("No", "Yes"),
                    levels = levels(arthritis$diabetes)),
  smoke = factor(c("No", "Yes"),
                 levels = levels(arthritis$smoke)))
```

```
)  
predict(fit, newdata = new_patients, type = "response")
```

```
1           2  
0.06001512 0.62751635
```